

# Computational Science on Many-Core Architectures:

## Exercise 8:

Viktor Beck, 11713110

December 2022

### Dot Product with OpenCL

1)

Here, we use the `vector_add.cpp` file as starting point for the dot product by simply converting the OpenCL kernel into a vector multiplication kernel ( $x[i] = x[i] * y[i]$ ). After the kernel call in the main function we sum the entries of vector `x` on the CPU - see Appendix. Note, that the code does only produce correct results when `reps = 1` (repetitions for measuring the execution time). But since the time is not dependent on the content of the vector it does not matter.

2) and 3)

For the performance comparison between CUDA and OpenCL an old CUDA code from exercise 2 was used. One can see that the CUDA implementation performs way better than the OpenCL implementation (... maybe because my CUDA dot product implementation is so good). Also, one can see the performance of the CPU which is clearly the worst. To get the CPU instead of the GPU running we simply have to exchange the 0 in the `platform_ids[0]` with 1 in the last line of the code below.

```
1 //  
2 // Query platform:  
3 //  
4 cl_uint num_platforms;  
5 cl_platform_id platform_ids[42]; //no more than 42 platforms supported...  
6 err = clGetPlatformIDs(42, platform_ids, &num_platforms); OPENCL_ERR_CHECK(err);  
7 std::cout << "# Platforms found: " << num_platforms << std::endl;  
8 cl_platform_id my_platform = platform_ids[0];
```

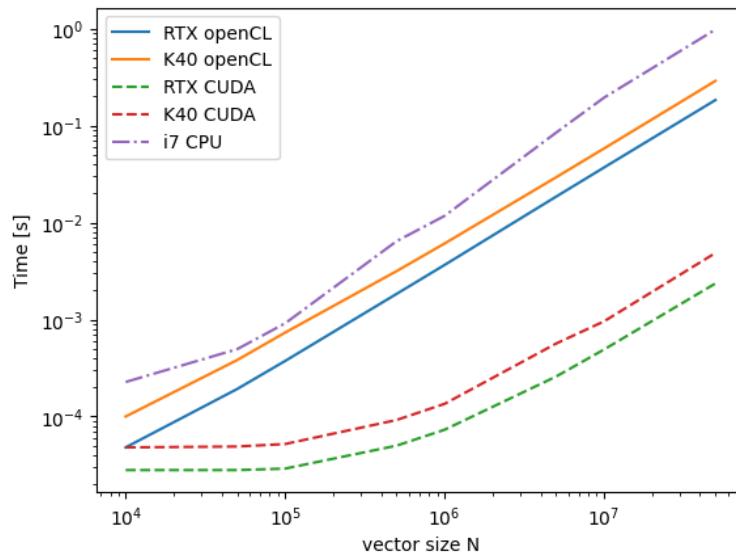


Figure 1: Execution times for different vector sizes of N for the dot product

## Appendix

Sorry for the horrible layout of the code listing here. The code was also uploaded with the report.

### Task 1),2),3)

```

1 //
2 // Tutorial for demonstrating a simple OpenCL vector addition kernel
3 //
4 // Author: Karl Rupp      rupp@iue.tuwien.ac.at
5 //
6
7 typedef double      ScalarType;
8
9
10 #include <iostream>
11 #include <string>
12 #include <vector>
13 #include <cmath>
14 #include <stdexcept>
15
16 #ifdef __APPLE__
17 #include <OpenCL/cl.h>
18 #else
19 #include <CL/cl.h>
20 #endif
21
22 // Helper include file for error checking
23 #include "ocl-error.hpp"
24 #include "timer.hpp"
25
26

```

```

27 const char *my_opengl_program = ""
28 "#pragma OPENCL EXTENSION cl_khr_fp64 : enable\n" // required to enable 'double' inside OpenGL pro
29 ""
30 "__kernel void dot_product(__global double *x,\n"
31 "                           __global double *y,\n"
32 "                           unsigned int N\n"
33 "{\n"
34 "    for (unsigned int i = get_global_id(0);\n"
35 "         i < N;\n"
36 "         i += get_global_size(0))\n"
37 "        x[i] = x[i] * y[i];\n"
38 "};"; // you can have multiple kernels within a single OpenGL program. For simplicity, this OpenGL pr
39
40
41 int main()
42 {
43     cl_int err;
44
45     //
46     ////////////////////////////////// Part 1: Set up an OpenGL context with one device //////////////////////////////////
47     //
48
49     //
50     // Query platform:
51     //
52     cl_uint num_platforms;
53     cl_platform_id platform_ids[42]; //no more than 42 platforms supported...
54     err = clGetPlatformIDs(42, platform_ids, &num_platforms); OPENCL_ERR_CHECK(err);
55     std::cout << "# Platforms found: " << num_platforms << std::endl;
56     cl_platform_id my_platform = platform_ids[0];
57
58
59     //
60     // Query devices:
61     //
62     cl_device_id device_ids[42];
63     cl_uint num_devices;
64     err = clGetDeviceIDs(my_platform, CL_DEVICE_TYPE_ALL, 42, device_ids, &num_devices); OPENCL_ERR_CHE
65     std::cout << "# Devices found: " << num_devices << std::endl;
66     cl_device_id my_device_id = device_ids[0];
67
68     char device_name[64];
69     size_t device_name_len = 0;
70     err = clGetDeviceInfo(my_device_id, CL_DEVICE_NAME, sizeof(char)*63, device_name, &device_name_len)
71     std::cout << "Using the following device: " << device_name << std::endl;
72
73     //
74     // Create context:
75     //
76     cl_context my_context = clCreateContext(0, 1, &my_device_id, NULL, NULL, &err); OPENCL_ERR_CHECK(er
77
78
79     //
80     // create a command queue for the device:
81     //
82     cl_command_queue my_queue = clCreateCommandQueueWithProperties(my_context, my_device_id, 0, &err);
83
84
85
86     //
87     ////////////////////////////////// Part 2: Create a program and extract kernels //////////////////////////////////
88     //

```

```

89
90     Timer timer;
91     timer.reset();
92
93     //
94     // Build the program:
95     //
96     size_t source_len = std::string(my_opengl_program).length();
97     cl_program prog = clCreateProgramWithSource(my_context, 1, &my_opengl_program, &source_len, &err);
98     err = clBuildProgram(prog, 0, NULL, NULL, NULL, NULL);
99
100    //
101    // Print compiler errors if there was a problem:
102    //
103    if (err != CL_SUCCESS) {
104
105        char *build_log;
106        size_t ret_val_size;
107        err = clGetProgramBuildInfo(prog, my_device_id, CL_PROGRAM_BUILD_LOG, 0, NULL, &ret_val_size);
108        build_log = (char *)malloc(sizeof(char) * (ret_val_size+1));
109        err = clGetProgramBuildInfo(prog, my_device_id, CL_PROGRAM_BUILD_LOG, ret_val_size, build_log, NULL);
110        build_log[ret_val_size] = '\0'; // terminate string
111        std::cout << "Log: " << build_log << std::endl;
112        free(build_log);
113        std::cout << "OpenCL program sources: " << std::endl << my_opengl_program << std::endl;
114        return EXIT_FAILURE;
115    }
116
117    //
118    // Extract the only kernel in the program:
119    //
120    cl_kernel my_kernel = clCreateKernel(prog, "dot_product", &err); OPENCL_ERR_CHECK(err);
121
122    std::cout << "Time to compile and create kernel: " << timer.get() << std::endl;
123
124
125    //
126    ////////////////////////////////////////////////// Part 3: Create memory buffers //////////////////////////////////////
127    //
128
129    //
130    // Set up buffers on host:
131    //
132    cl_uint vector_size = 50000000;
133    std::vector<ScalarType> x(vector_size, 3.0);
134    std::vector<ScalarType> y(vector_size, 2.0);
135
136    std::cout << std::endl;
137    std::cout << "Vectors before kernel launch:" << std::endl;
138    std::cout << "x: " << x[0] << " " << x[1] << " " << x[2] << " ..." << std::endl;
139    std::cout << "y: " << y[0] << " " << y[1] << " " << y[2] << " ..." << std::endl;
140
141    //
142    // Now set up OpenCL buffers:
143    //
144    cl_mem ocl_x = clCreateBuffer(my_context, CL_MEM_READ_WRITE | CL_MEM_COPY_HOST_PTR, vector_size * sizeof(ScalarType), x.data(), &err);
145    cl_mem ocl_y = clCreateBuffer(my_context, CL_MEM_READ_WRITE | CL_MEM_COPY_HOST_PTR, vector_size * sizeof(ScalarType), y.data(), &err);
146
147
148    //
149    ////////////////////////////////////////////////// Part 4: Run kernel //////////////////////////////////////
150    //

```

```

151     size_t local_size = 128;
152     size_t global_size = 128*128;
153
154
155     //
156     // Set kernel arguments:
157     //
158     err = clSetKernelArg(my_kernel, 0, sizeof(cl_mem), (void*)&ocl_x); OPENCL_ERR_CHECK(err);
159     err = clSetKernelArg(my_kernel, 1, sizeof(cl_mem), (void*)&ocl_y); OPENCL_ERR_CHECK(err);
160     err = clSetKernelArg(my_kernel, 2, sizeof(cl_uint), (void*)&vector_size); OPENCL_ERR_CHECK(err);
161
162     std::vector<double> times;
163     int reps = 6;
164     double sum;
165
166     // START TIMING HERE:
167
168     for (int i=0; i<reps; i++){
169         timer.reset();
170
171         //
172         // Enqueue kernel in command queue:
173         //
174         err = clEnqueueNDRangeKernel(my_queue, my_kernel, 1, NULL, &global_size, &local_size, 0, NULL, NU
175
176         // wait for all operations in queue to finish:
177         err = clFinish(my_queue); OPENCL_ERR_CHECK(err);
178
179
180         //
181         ////////////////////////////////// Part 5: Get data from OpenCL buffer //////////////////////////////////
182         //
183
184         err = clEnqueueReadBuffer(my_queue, ocl_x, CL_TRUE, 0, sizeof(ScalarType) * x.size(), &x[0], 0,
185
186         // summing like a pro on the CPU
187         sum = 0;
188         for (int i = 0; i < vector_size; i++){
189             sum += x[i];
190         }
191         times.push_back(timer.get());
192     }
193
194     std::cout << "Exec. time\n" << times[reps/2] << std::endl;
195
196     // TIMER END
197
198     std::cout << std::endl;
199     std::cout << "Vectors after kernel execution:" << std::endl;
200     std::cout << "x: " << x[0] << " " << x[1] << " " << x[2] << " ..." << std::endl;
201     std::cout << "y: " << y[0] << " " << y[1] << " " << y[2] << " ..." << std::endl;
202
203
204     std::cout << "Dot product of x and y = " << sum << std::endl;
205
206
207     //
208     // cleanup
209     //
210     clReleaseMemObject(ocl_x);
211     clReleaseMemObject(ocl_y);
212     clReleaseProgram(prog);

```

```
213     clReleaseCommandQueue(my_queue);
214     clReleaseContext(my_context);
215
216     std::cout << std::endl;
217     std::cout << "#" << std::endl;
218     std::cout << "# My first OpenCL application finished successfully!" << std::endl;
219     std::cout << "#" << std::endl;
220     return EXIT_SUCCESS;
221 }
```